



SELA | DEVELOPER | PRACTICE
July 3-5, 2018

Kevin Gosse @kookiz
Grégory Léocadie @gleocadie
Christophe Nasarre @chnasarre

criteo.

.NET Core Monitoring









Criteo in numbers

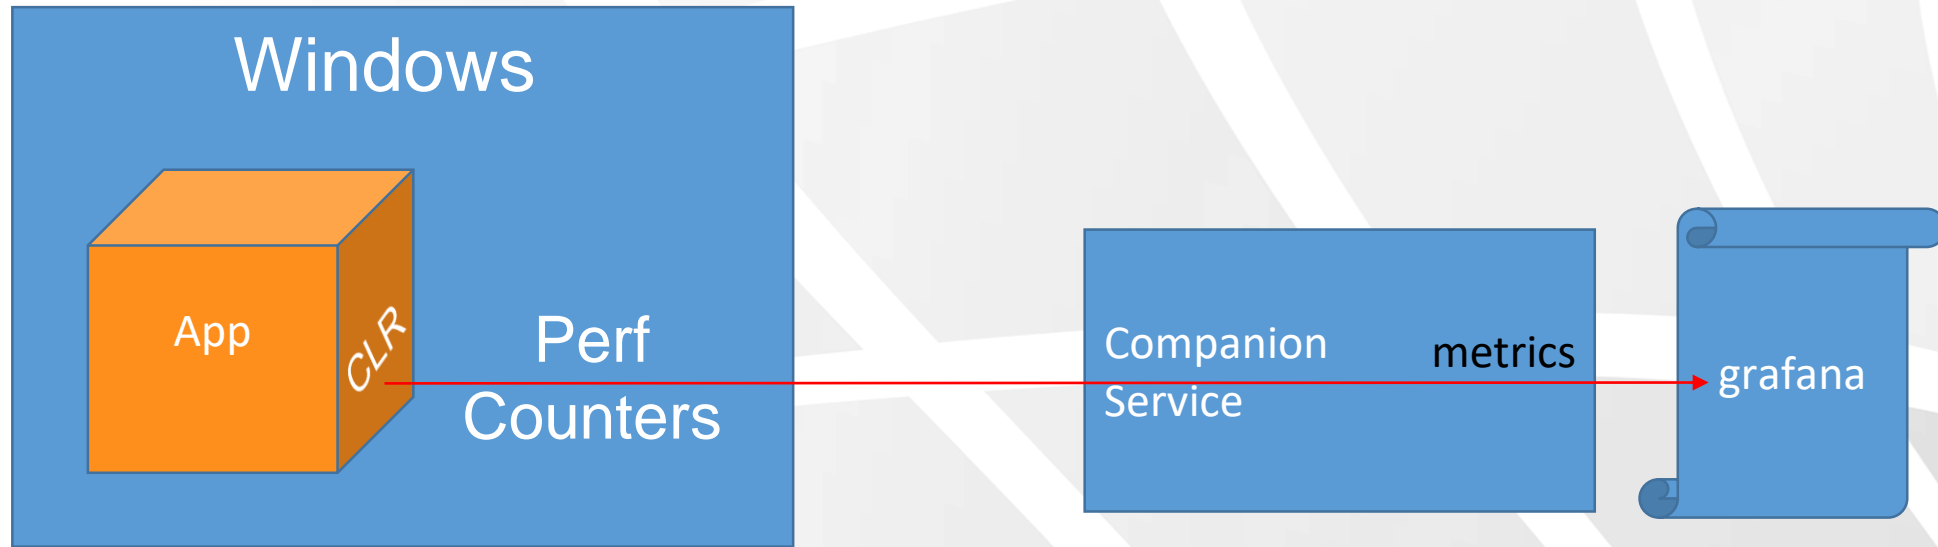
- ⚡ 9000+ servers running Windows
- ⚡ 4000+ front-end servers
- ⚡ 200+ billions requests per day



Currently on .NET 4.6.2 and... moving to .NET Core and Linux



Performance counters pipeline to Grafana



Performance counters limitations

✦ Misleading counters

- ✦ Gen0 size
- ✦ Gen collections count
- ✦ Threads count

✦ Missing information

- ✦ Suspension/contention time
- ✦ Exceptions/finalizers

✦ **No perf counters in .NET Core**

- ✦ Better than being blind...
-

Alternative to Perf Counters: CLR traces

- ✦ On Windows, Perfview has been using ETW for years
 - ✦ Send to ETW on Windows and to LTTng on Linux
 - ✦ Stubs generated in CLR code at compile time via scripts\Python files
 - ✦ Shared method that produces traces
 - ✦ Look for **FireEtw<xxx>** and **ETW::<type>::<methods>** functions
 - ✦ Great to find the undocumented events ;^)
-

Monitor Memory

✦ **Garbage Collections**

✦ *GCHeapStats*

- ✦ All generations size + pinned objects + promoted sizes

✦ **Application Threads Suspension Time during GC**

✦ *GCSuspendEEStart/GCRestartEEStop*

✦ **Sampling heap allocations**

✦ *AllocationTick*

- ✦ Type name
 - ✦ Allocation size since the last event
 - ✦ LOH/SOH
-

Monitor Threading

✦ **Threads**

- ✦ *ThreadCreating/ThreadRunning*

- ✦ No way to know when a thread ends!

✦ **ThreadPool Worker threads**

- ✦ *ThreadPoolWorkerThreadStart / ThreadPoolWorkerThreadStop*

- ✦ Active/retired count

✦ **ThreadPool I/O threads**

- ✦ *IOThreadCreationStart / IOThreadCreationStop*

- ✦ Active/retired count

Monitoring/Investigation bonus!

✧ **Exceptions Thrown**

✧ *ExceptionStart*

- Exception type
- Exception message

✧ **Finalizers**

✧ *GCFinalizeObject/GCFinalizersStart/GCFinalizersEnd*

- Type name of the finalized instance
- Count and time spent

✧ **Contention time**

✧ *ContentionStart/ContentionStop*



Connecting to CLR with TraceEvent (Windows)

✦ Creating a `TraceEventSession`

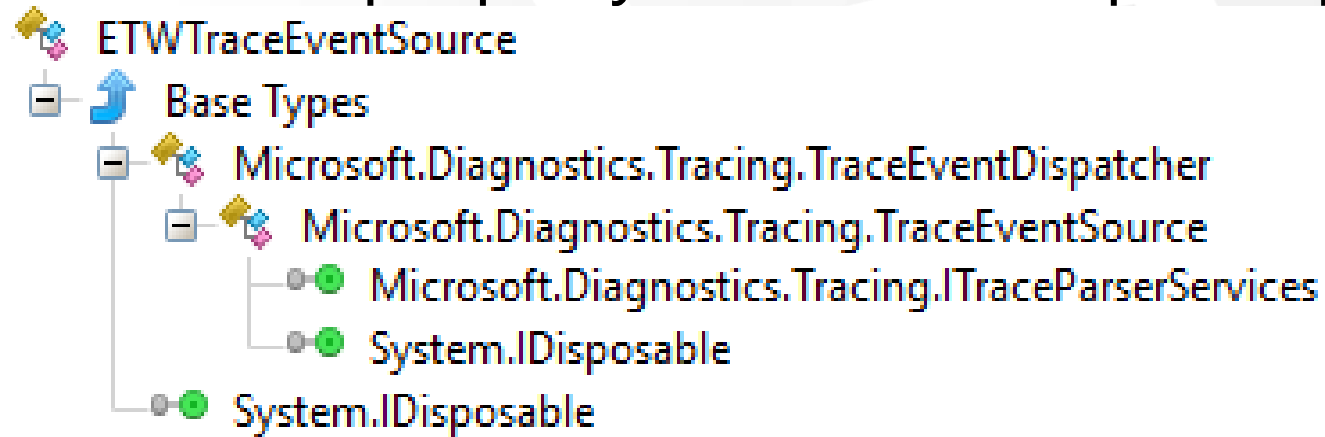
- ✦ Give it a unique name
- ✦ Don't forget to dispose it
 - ✦ `logman -ets stop <session name>`

✦ Enable `ClrTraceEventParser.ProviderGuid`

- ✦ Pick the [right event level](#) (depending on [events](#))
 - ✦ Filter events by [keywords](#)
 - ✦ Additional options set via `TraceEventProviderOptions`
-

Decyphering CLR traces with TraceEvent

- ✦ The source property on session exposes parsers



- ✦ Use the parser exposed by the source Clr property
 - ✦ The All event is raised each time a trace is received
 - ✦ Start the processing of events thanks to Process()
 - ✦ This is a blocking call!
-

Naive usage of TraceEvent

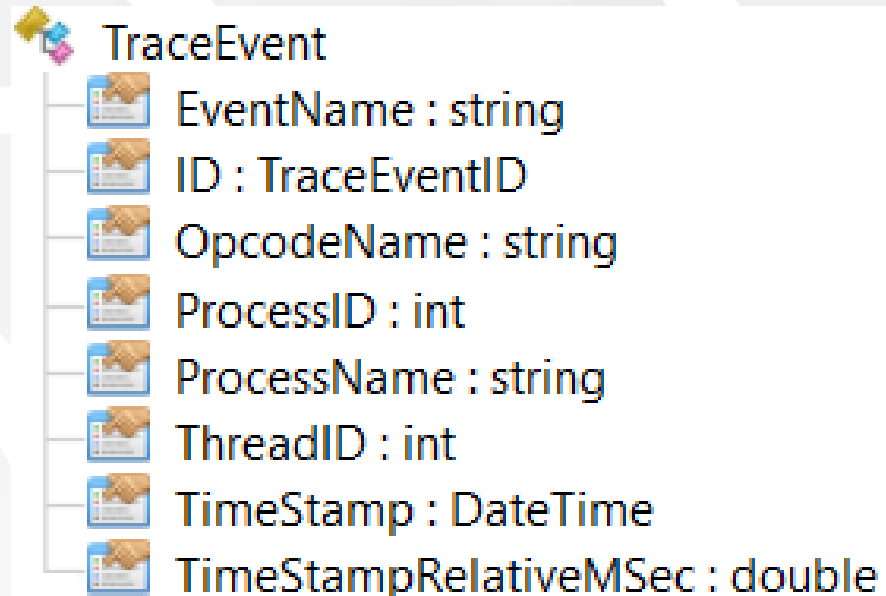
Demo

- Notion of session
- Setup providers
- Listen to All events



Decyphering CLR events with TraceEvent

- ✦ Listen to CLR parser [events](#)
 - ✦ Dig into the documentation for details
- ✦ Each event has a common **TraceEvent** payload



Decyphering CLR events with TraceEvent

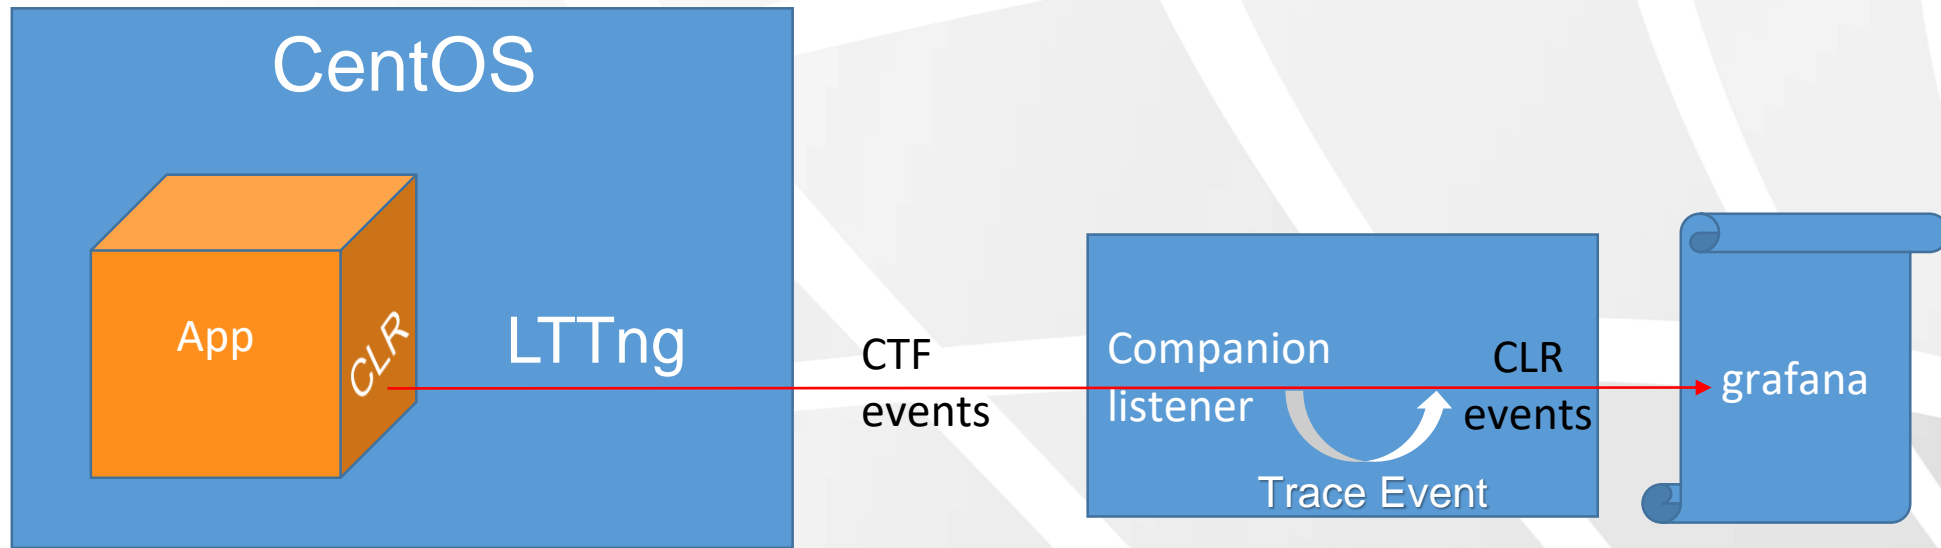
Demo

- First chance exception: unique event
- Thread contention: start/stop events





CentOS CLR events pipeline to Grafana



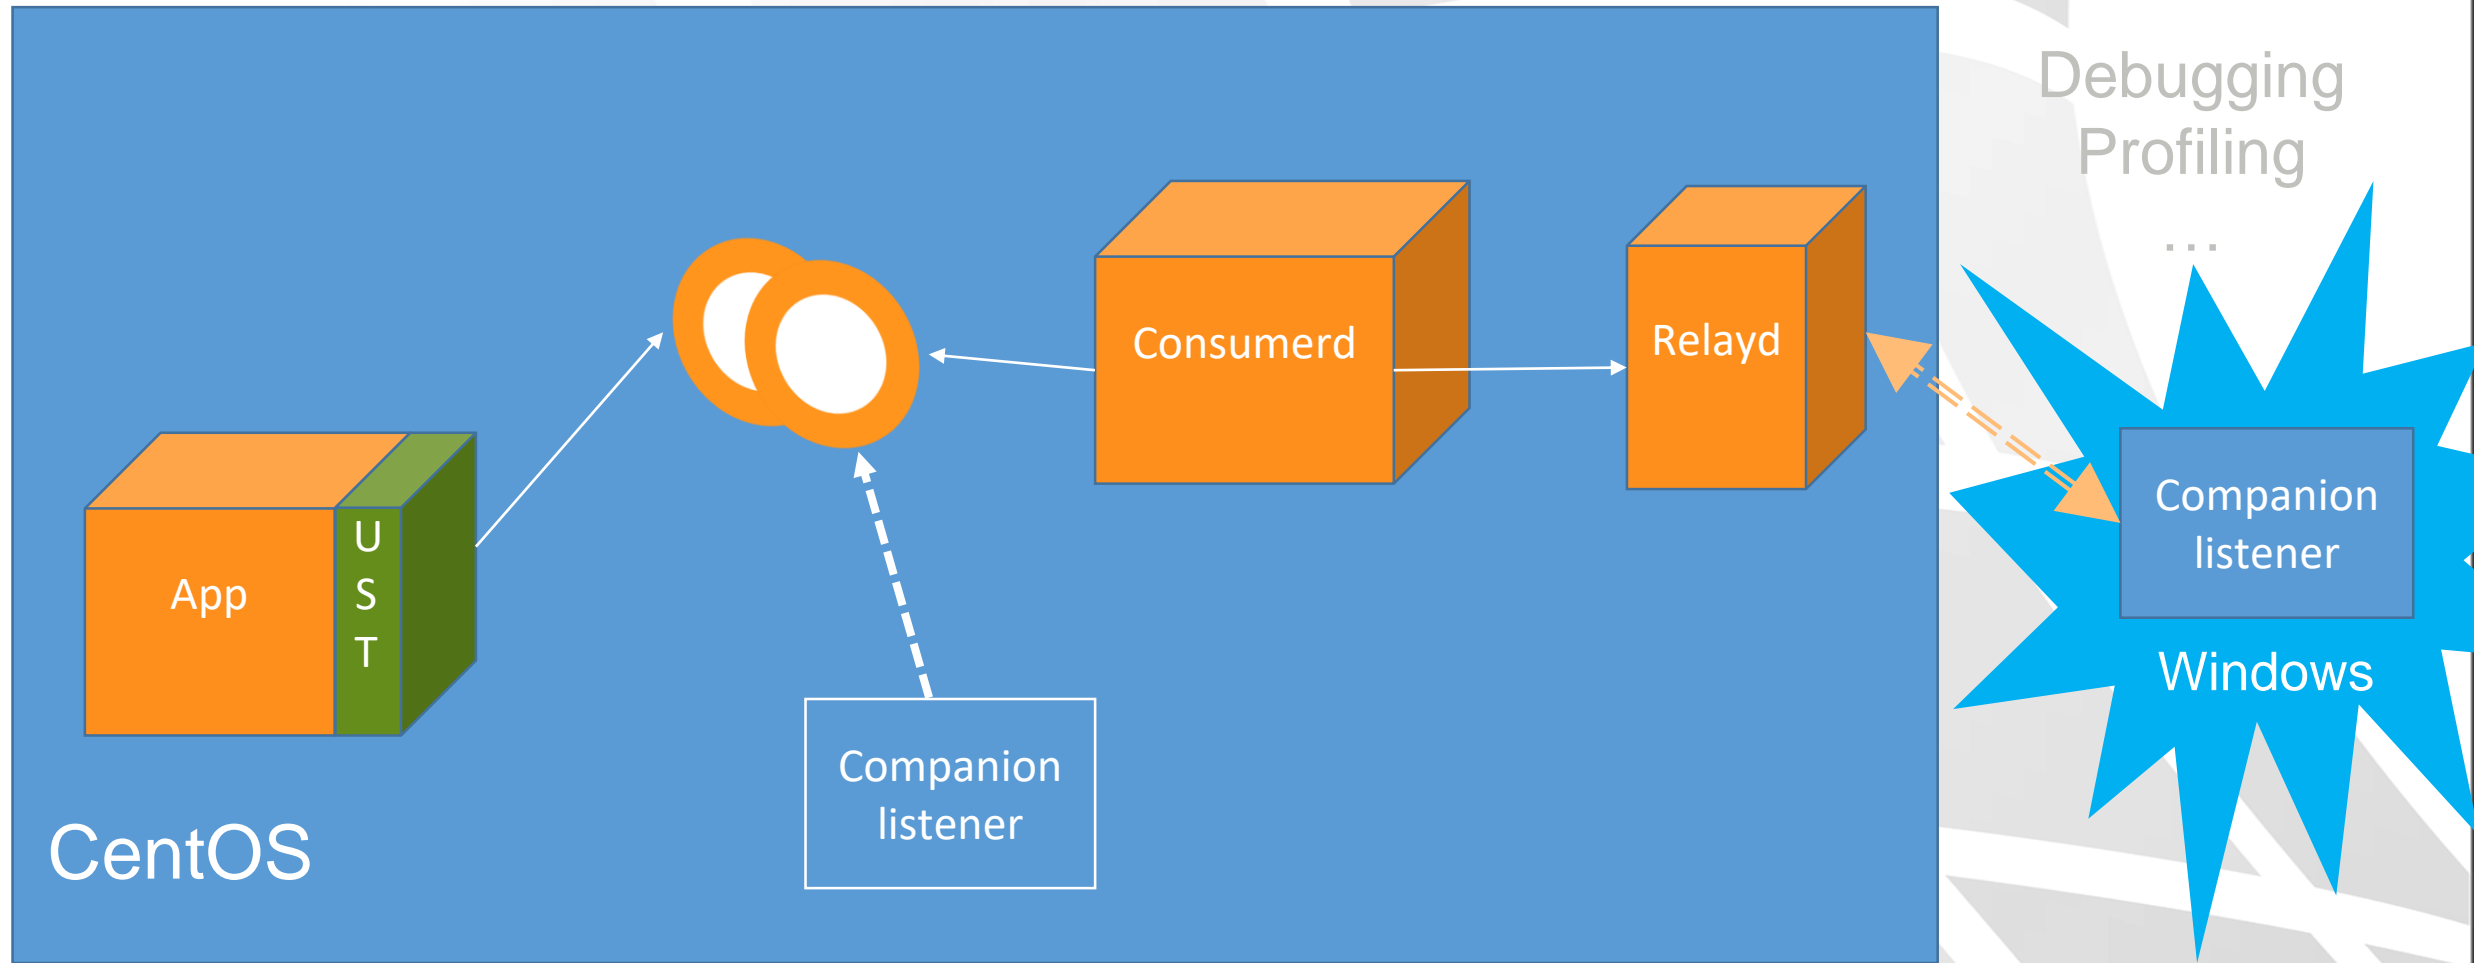
Decyphering CLR event produced as CTF traces

- ✦ Use TraceEvent library but... only file-based ctor

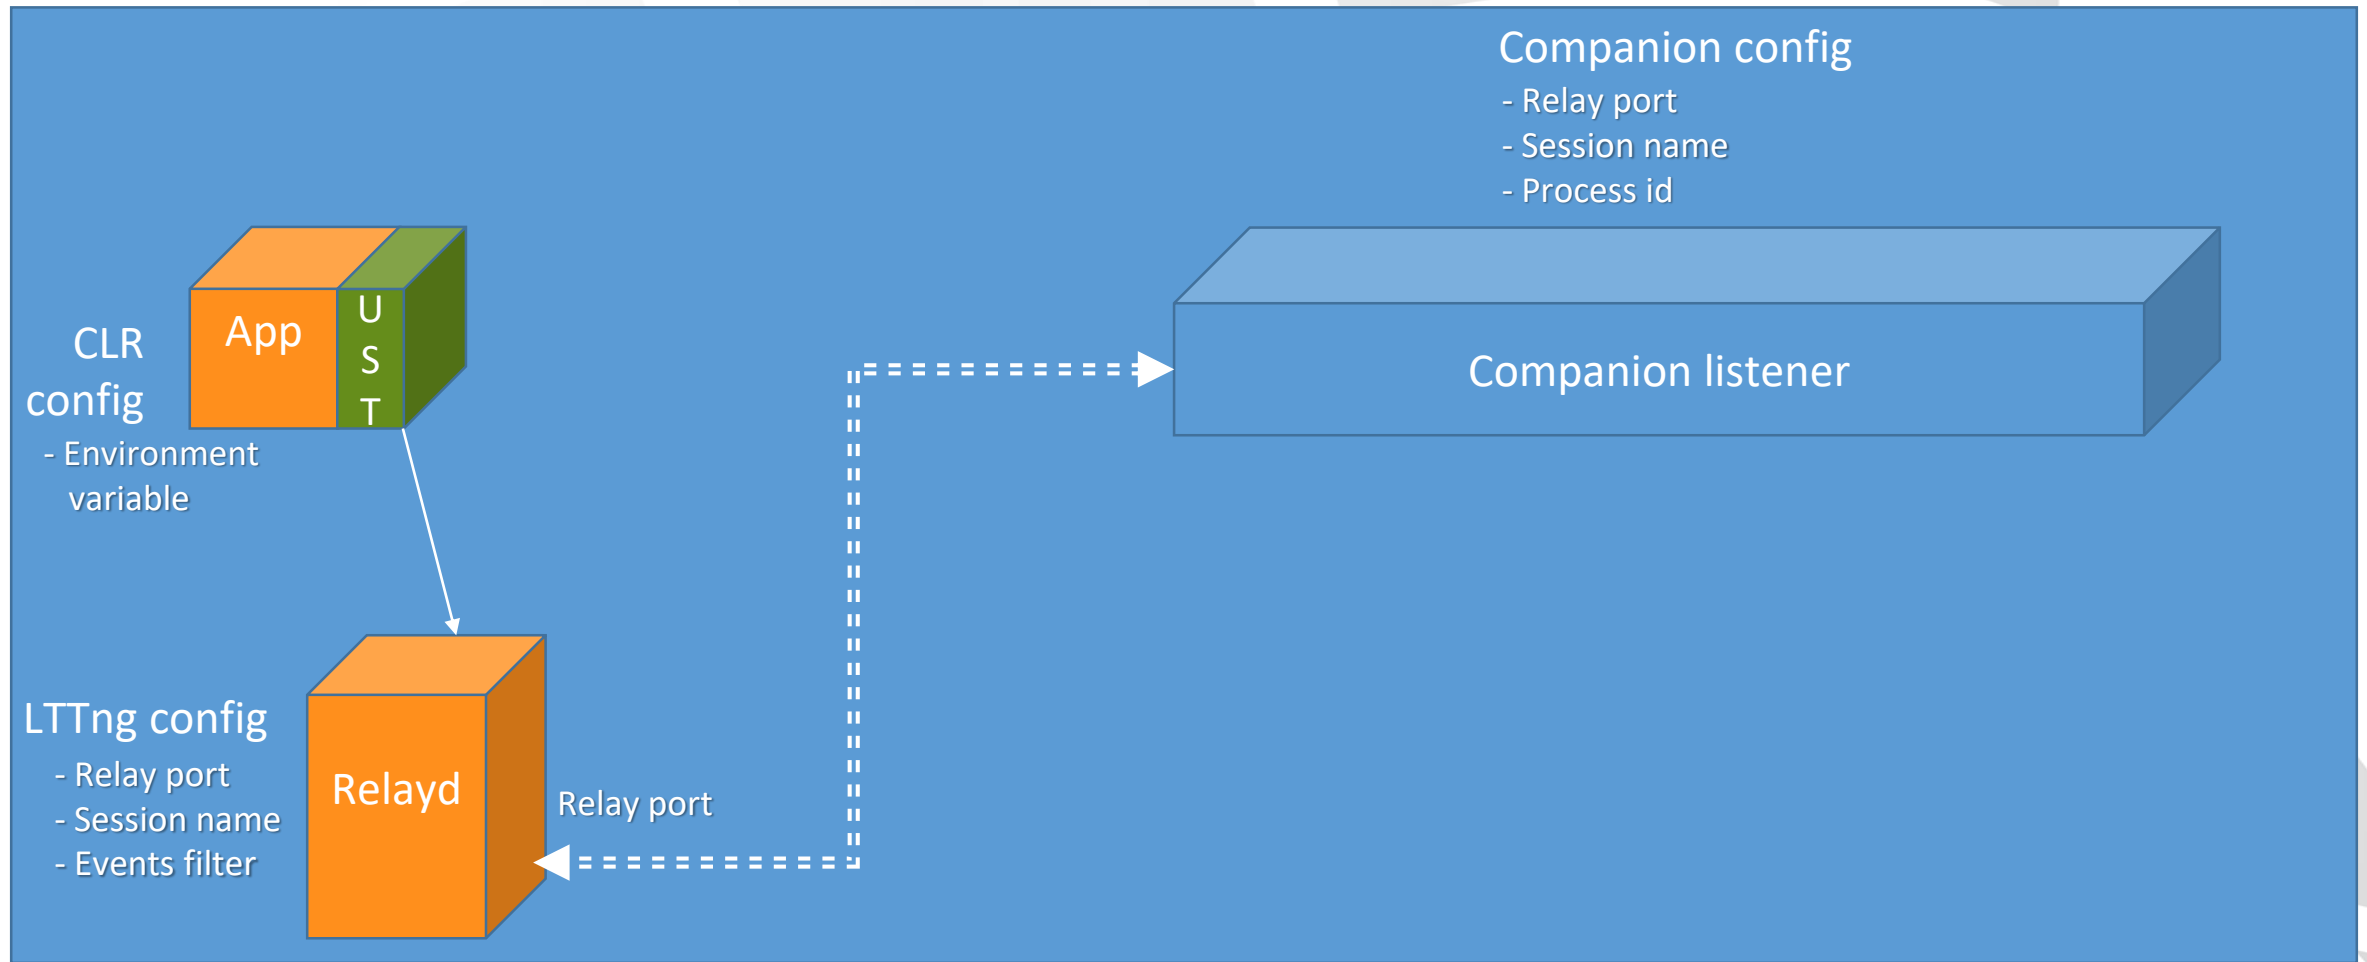
```
44  
45     public CtfTraceEventSource(string fileName)  
46     {
```

- ✦ Implement CTF live session support:
<https://github.com/Microsoft/perfview/pull/340>
-

How to integrate the LTTng pipeline in Criteo



LTNg Live Session implementation details



Configuring pipeline on Linux

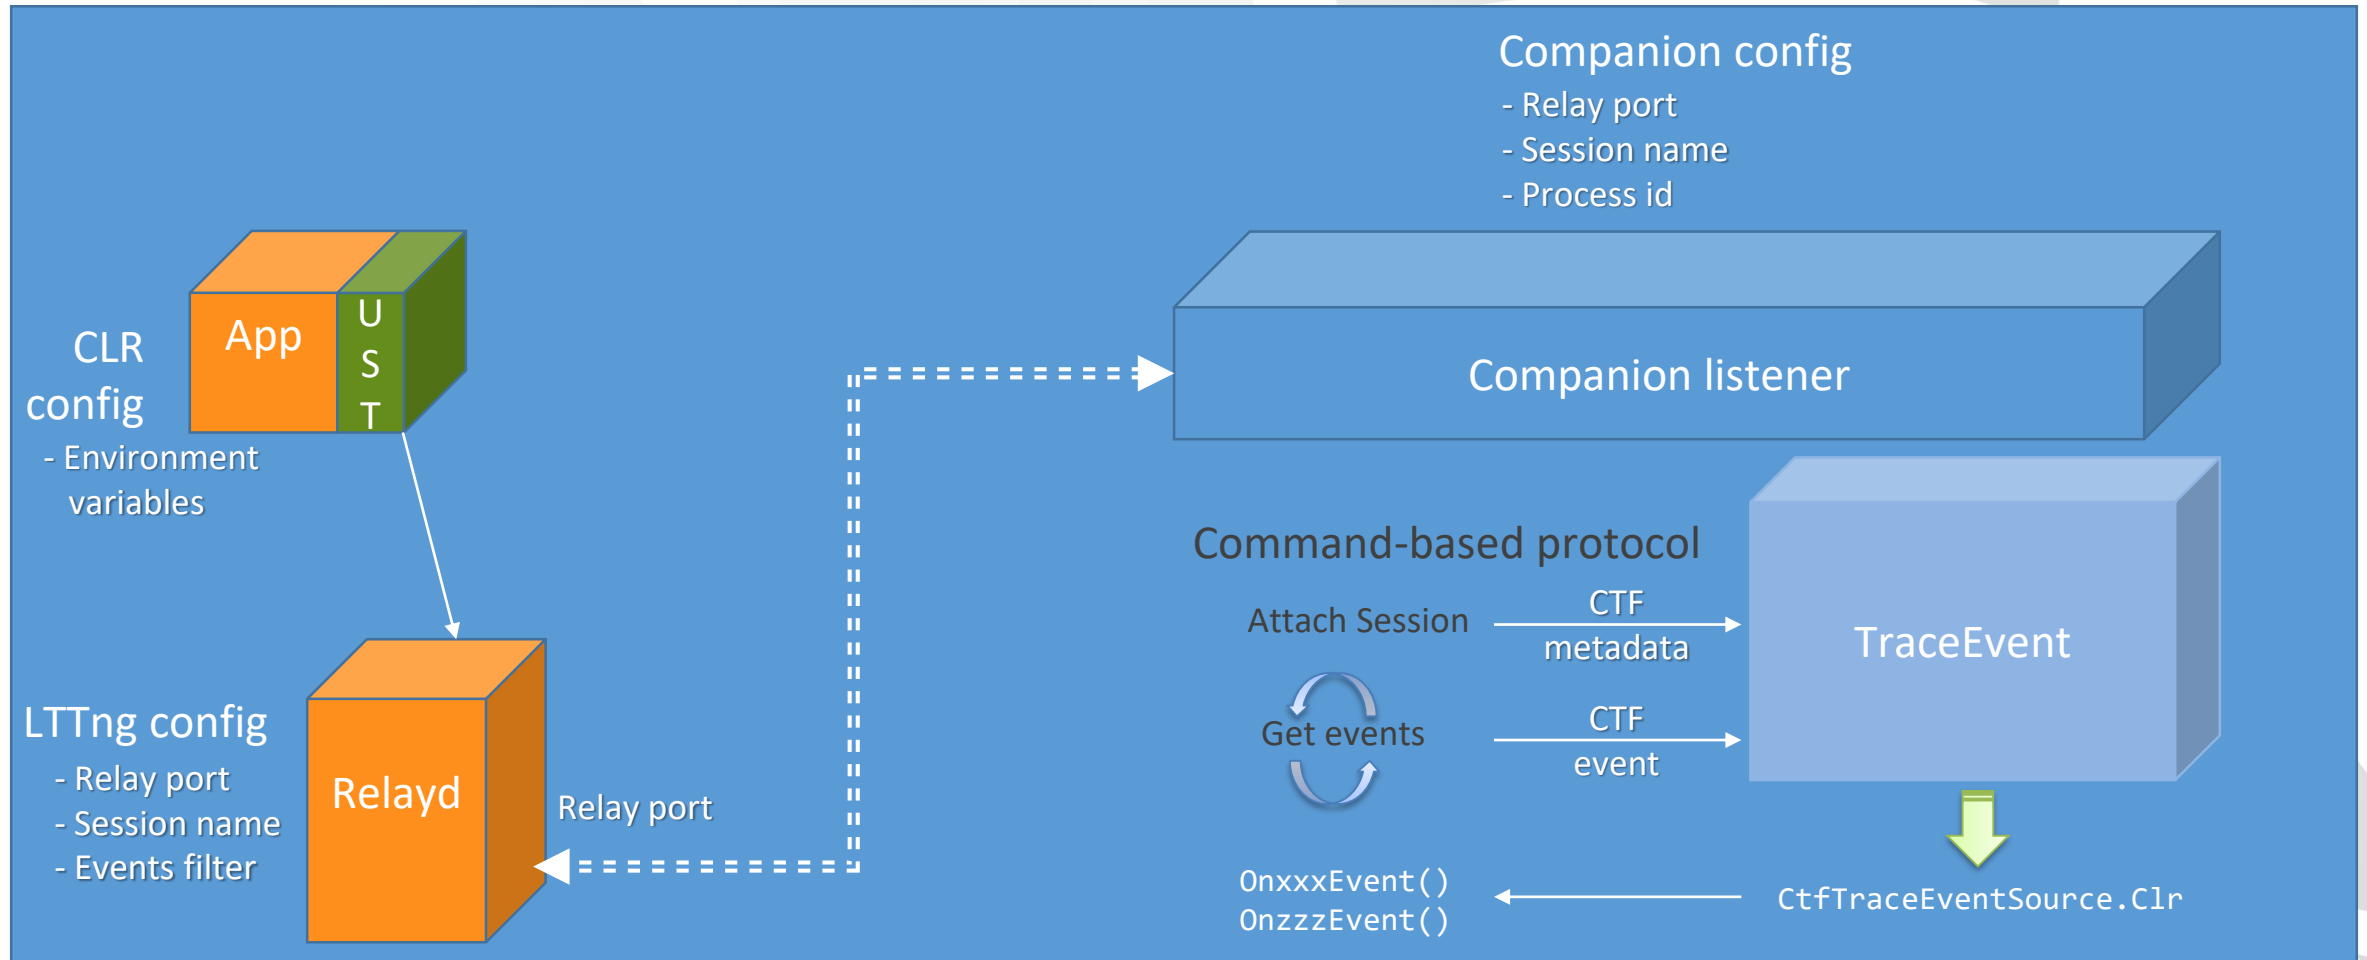
✦ LTTng configuration


```
lttng-relayd -d --control-port=tcp://0.0.0.0:8080  
--data-port=tcp://0.0.0.0:8081 --live-port=tcp://0.0.0.0:8088  
  
lttng create sela_conference_session --live=100000  
--ctrl-url=tcp://0.0.0.0:8080 --data-url=tcp://0.0.0.0:8081  
  
lttng enable-event --userspace -c sela_conference_session  
--tracepoint DotNETRuntime:*
```

✦ Environment variables for CLR

```
✦ export COMPlus_PerfMapEnabled=1  
✦ export COMPlus_EnableEventLog=1
```

LTtng Live Session implementation details





D:\LTTngDemo>EventTracing.Linux.Examples.exe

LTNg pitfalls

- ✦ The relay is also storing events on disk
 - ✦ Customizable by log rotation
 - ✦ LTNg seems to have a lower throughput than ETW
 - ✦ Filter your CLR events carefully in LTNg configuration (ex: AllocationTick)
-

Key takeaways

- ✦ Prefer CLR traces to performance counters
- ✦ No alternative when moving to .NET Core
- ✦ Learning curve moving to Linux could be...



Blog series about performance counters and ETW:

<http://labs.criteo.com/2018/06/replace-net-performance-counters-by-clr-event-tracing/>

Questions

